

Projecting chains of constrained points

Morten Silcowitz

Feb 2026

1 Formulation

We're going to solve

$$\begin{aligned} \min_{\mathbf{x}} \quad & \frac{1}{2} (\mathbf{x} - \mathbf{p})^T \mathbf{M} (\mathbf{x} - \mathbf{p}) \\ \text{s.t.} \quad & \mathbf{z} = \mathbf{R}\mathbf{x} \quad \text{and} \quad \mathbf{z}_i^T \mathbf{z}_i = 1 \end{aligned} \tag{1}$$

where

$n \in \mathbb{N}$	is the number of particles between links
$m \in \mathbb{N}$	is the $n - 1$ of links between particles
$\mathbf{x} \in \mathbb{R}^{3n}$	is the solution vector,
$\mathbf{p} \in \mathbb{R}^{3n}$	is target/free moving positions,
$\mathbf{z} \in \mathbb{R}^{3m}$	is the vector of difference vectors in the chain,
$\mathbf{z}_i \in \mathbb{R}^3$	is the i th 3-vector in \mathbf{z} ,
$\mathbf{R} \in \mathbb{R}^{3m \times 3n}$	is the per vector finite-difference matrix so that $(\mathbf{R}\mathbf{x})_i = \mathbf{x}_i - \mathbf{x}_{i+1}$
$\mathbf{M} \in \mathbb{R}^{3n \times 3n}$	is the mass matrix, simply having the mass (or weight) of each particle in each 3x3 block diagonal

We want to be working only on the difference vectors in \mathbf{z} . To do that, we need to eliminate \mathbf{x} , and we do this by holding \mathbf{z} constant and deriving the optimality conditions for (1). We first state the Lagrangian:

$$\mathcal{L}(\mathbf{x}, \lambda) = \frac{1}{2} (\mathbf{x} - \mathbf{p})^T \mathbf{M} (\mathbf{x} - \mathbf{p}) + \lambda^T (\mathbf{z} - \mathbf{R}\mathbf{x})$$

with the optimality conditions:

$$\begin{aligned} \mathbf{M}(\mathbf{x} - \mathbf{p}) - \mathbf{R}^T \lambda &= 0 \\ \mathbf{z} - \mathbf{R}\mathbf{x} &= 0 \end{aligned}$$

solving for λ gives us:

$$\lambda = \underbrace{(\mathbf{RM}^{-1}\mathbf{R}^T)}_S^{-1}(\mathbf{z} - \mathbf{Rp})$$

where $\mathbf{S} = \mathbf{RM}^{-1}\mathbf{R}^T \in \mathbb{R}^{3m \times 3m}$ is a block tridiagonal SPD matrix. We use this to give us an expression for \mathbf{x} :

$$\mathbf{x} = \mathbf{p} + \mathbf{M}^{-1}\mathbf{R}^T\mathbf{S}^{-1}(\mathbf{z} - \mathbf{Rp})$$

Inserting this into (1) eliminates $\mathbf{z} = \mathbf{Rx}$ and we are left with

$$\begin{aligned} \min_{\mathbf{x}} \quad & \frac{1}{2} (\mathbf{z} - \mathbf{Rp})^T \mathbf{S}^{-1} (\mathbf{z} - \mathbf{Rp}) \\ \text{s.t.} \quad & \mathbf{z}_i^T \mathbf{z}_i = 1 \quad i = 1 \dots m \end{aligned} \tag{2}$$

In this reformulated problem we must find a \mathbf{z} that consist of unit length 3-vectors that minimize the distance to \mathbf{Rp} under the \mathbf{S}^{-1} -norm. Note that it is trivial to project \mathbf{z} to satisfy the unit length constraints (under the regular \mathbf{I} -norm, known as normalization), and likewise it is trivial to obtain \mathbf{x} given any \mathbf{z} . This means that whenever we take a newton step on \mathbf{z} we can trivially reign it in to its feasible solution, and in effect only change the directions of each \mathbf{z}_i vector, never their lengths.

2 Solution

We want to solve this system using a newton-like method, so to make sense of that we first derive the full newton step. The Lagrangian of (2) is

$$\mathcal{L}_{(\mathbf{z}, \lambda)} = \frac{1}{2}(\mathbf{z} - \mathbf{Rp})^T \mathbf{S}^{-1} (\mathbf{z} - \mathbf{Rp}) + \frac{1}{2} \sum_{i=1}^m \lambda_i (\mathbf{z}_i^T \mathbf{z}_i - 1) \tag{3}$$

The optimality conditions for (3) are:

$$\begin{aligned}\mathbf{S}^{-1}(\mathbf{z} - \mathbf{R}\mathbf{p}) + \mathbf{Q}_z^\top \lambda &= 0 \\ \mathbf{z}_i^\top \mathbf{z}_i - 1 &= 0 \quad i = 1, \dots, m\end{aligned}$$

and the newton step to find a stationary point (\mathbf{z}, λ) for (3) is

$$\underbrace{\begin{bmatrix} \mathbf{S}^{-1} + \mathbf{D}_\lambda & \mathbf{Q}_z^\top \\ \mathbf{Q}_z & \mathbf{0} \end{bmatrix}}_{H\mathcal{L}_{(\mathbf{z}, \lambda)}} \begin{bmatrix} \Delta \mathbf{z} \\ \Delta \lambda \end{bmatrix} = \underbrace{\begin{bmatrix} \mathbf{S}^{-1}(\mathbf{z} - \mathbf{R}\mathbf{p}) + \mathbf{Q}_z^\top \lambda \\ \mathbf{Q}_z \mathbf{z} - \mathbf{1} \end{bmatrix}}_{\nabla \mathcal{L}_{(\mathbf{z}, \lambda)}} \quad (4)$$

where

$$\mathbf{Q}_z = \begin{bmatrix} \mathbf{z}_1^\top & 0 & \cdots & 0 \\ 0 & \mathbf{z}_2^\top & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \mathbf{z}_m^\top \end{bmatrix} \in \mathbb{R}^{m \times 3m} \quad \mathbf{D}_\lambda = \begin{bmatrix} \mathbf{I}_{\lambda_1} & 0 & \cdots & 0 \\ 0 & \mathbf{I}_{\lambda_2} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \mathbf{I}_{\lambda_m} \end{bmatrix} \in \mathbb{R}^{3m \times 3m}$$

There are many interesting ways to go about solving this system, but in our case the goal is to exploit the tri-diagonal structure of \mathbf{S} , in order to use fast direct solves. Instead of solving the full step, we first solve just for $\Delta \lambda$ at $\lambda = \mathbf{0}$, effectively resetting the multipliers in each iteration. After that we solve for Δz to update the solution variables.

2.1 Solving for λ

To derive the solution for $\Delta \lambda$ we eliminate Δz by using $\lambda = \mathbf{0}$ and solving the first row in (4). Interestingly, because we will project the vectors in \mathbf{z} in each iteration the $\mathbf{Q}_z \mathbf{z} - \mathbf{1}$ term is always 0, which will simplify things a bit:

$$\Delta \mathbf{z} = \mathbf{S}(\mathbf{Q}_z^\top \Delta \lambda + \mathbf{S}^{-1}(\mathbf{z} - \mathbf{R}\mathbf{p}))$$

Substitute Δz into the second row in (4):

$$Q_z(S(Q_z^T \Delta \lambda + S^{-1}(z - Rp))) = \underbrace{Q_z z - 1}_0$$

Simplify to solve for $\Delta \lambda$ we get:

$$Q_z S Q_z^T \Delta \lambda = -Q_z(z - Rp)$$

This system is tridiagonal because the operator $Q_z(\cdot)Q_z^T$ collapses each 3×3 block in S to a scalar value $z_i^T S_{(i,j)} z_j$. This operation effectively reduces the block tridiagonal structure of S to the plain tridiagonal structure in $Q_z S Q_z^T \in \mathbb{R}^{m \times m}$. Also note that $Q_z z = 1$ appears in the RHS and so it could be further reduced to $Q_z R_p - 1$ but since the vector $z - Rp$ is needed elsewhere this has neglectable practical benefit.

2.2 Solving for Δz

To solve for Δz , we substitute $\Delta \lambda = 0$ into the first row of (4):

$$(S^{-1} + D_\lambda) \Delta z = S^{-1}(z - Rp) - Q_z^T \lambda \quad (5)$$

To avoid having to solve the dense system $S^{-1} + D_\lambda$ we exploit that we can cheaply compute the cholesky decomposition $LL^T = S$ and make use of the identity

$$S^{-1} + D_\lambda = L^{-T}(I + L^T D_\lambda L)L^{-1} \quad (6)$$

we insert (6) into (5)

$$L^{-T}(I + L^T D_\lambda L)L^{-1} \Delta z = L^{-T} L^{-1}(z - Rp) - Q_z^T \lambda$$

then multiplying with the invertible \mathbf{L}^\top

$$(\mathbf{I} + \mathbf{L}^\top \mathbf{D}_\lambda \mathbf{L}) \mathbf{L}^{-1} \Delta \mathbf{z} = \mathbf{L}^{-1} (\mathbf{z} - \mathbf{R}\mathbf{p}) - \mathbf{L}^\top \mathbf{Q}_z^\top \lambda$$

means we can solve for $\mathbf{u} = \mathbf{L}^{-\top} \Delta z$ and then obtain Δz like so

$$(\mathbf{I} + \mathbf{L}^\top \mathbf{D}_\lambda \mathbf{L}) \mathbf{u} = \mathbf{L}^{-1} (\mathbf{z} - \mathbf{R}\mathbf{p}) - \mathbf{L}^\top \mathbf{Q}_z^\top \lambda$$

and then finally

$$\Delta \mathbf{z} = \mathbf{L}^\top \mathbf{u}$$

The great advantage here is that $(\mathbf{I} + \mathbf{L}^\top \mathbf{D}_\lambda \mathbf{L})$ remains block tridiagonal and so it's fast to solve. The disadvantage is that we need one extra solve of the block bidiagonal \mathbf{L} to compute $\mathbf{L}^{-1}(\mathbf{z} - \mathbf{R}\mathbf{p})$. This is annoying but still greatly outweighs the pain of solving a dense system using iterative methods. Note also that in the trivial case where $\lambda = 0$ the system reduces to the identity. In an effort to keep the eigenvalues of $(\mathbf{I} + \mathbf{L}^\top \mathbf{D}_\lambda \mathbf{L})$ positive, we clamp λ to positive values before we compute \mathbf{D}_λ . This is a hack but seems to work well in practice. We now proceed to list the full algorithm.

3 Algorithm

Here we'll outline the algorithm in full before explaining each step in detail:

```

01 L = cholesky(S)
02 z = Rx
03 while true do
04     Q, z = normalize(z)
05     y = Rp - z
06     λ = solve(QSQT, Qy)
```

```

07     bz = solve(L, y)
08     bλ = LTQTλ
09     if ||bz - bl||2 < ε then
10         s = solve(LT, bz)
11         x = p - M-1RTs
12         return x
13     D = diagonal(max(0, λ))
14     Δz = L solve(I + LTDL, bz - bλ)
15     z = z + Δz

```

Each iteration the algorithm does these main steps:

1. Normalize the 3-vectors in \mathbf{z}
2. Compute the $n - 1$ Lagrange multipliers of the system at \mathbf{z}
3. Check the solution residual, if below threshold compute and return \mathbf{x}
4. Compute the Newton step on \mathbf{z} and go to step 1

Note the following:

- we compute the cholesky decomposition $\mathbf{L}\mathbf{L}^T = \mathbf{S}$ to keep the systems we need to solve sparse and symmetric. See more details bellow. Computing cholesky for a block-tridiagonal SPD matrix is simple and fast, and if \mathbf{M} is fixed it can be done off-line/once.
- every `solve()` is at least a block tri-diagonal system (or simpler), which are fast and easy to solve
- to compute the Lagrange multipliers we build the Jacobian matrix $\mathbf{Q} \in \mathbb{R}^{3m \times m}$, which is simply the \mathbf{z}_i vectors vertically stacked

- the matrix \mathbf{D} is the contribution to the Hessian stemming from the non-linear unit length constraints. We clamp it to be positive to keep the Hessian PD. This is a hack but seems to play out well in practice